Notes on Solace Veil V21 · Article Mapping (refined)

Reading against Annex A v1.0 · External Reviewer observation 12 May 2026

To:	Tim Zlomke, Managing Partner Neurovia Dynamics · Founder Moral Clarity AI · builder Solace Veil
From:	Lamar B. Shucrani, Founder, Sprinkling Act, Brussels
Date:	13 May 2026
Reference:	Shucrani, L. B. (2026). EU AI Act Readiness Report (April 2026) + Annex A v1.0 (May 2026). Zenodo. DOI 10.5281/zenodo.19671329 · 10.5281/zenodo.20042175

> External Reviewer access to Solace Veil V21 (session SV-7X91-A77C, 12 May 2026) confirms what the May 9 prototype anticipated and refines two points: the runtime vocabulary maps onto article-level obligations at finer resolution than the synthetic-scenario reading allowed, and the multi-deployer coherence problem under independently compliant deployer obligations is the load-bearing edge where the regulatory layer does not yet name the architecture. The mapping below extends Annex A §7.1 base cascade (§2 + §5 + §6) with the runtime vocabulary directly.

1. CONVERGENCE AT FINER RESOLUTION

Three runtime concepts whose article-level correspondence is closer than synthetic-scenario inference suggested.

— **T=0 state (Execution Kernel)** maps onto **Article 26§5** deployer monitoring and immediate-suspension obligation. The runtime treats T=0 as the binding moment of authority + revocation checks; the Act treats §5 as the binding obligation when there is **"reason to consider"** a §79(1) risk. Same architecture under different vocabulary.

— **Revalidation queues (Reality Model)** map onto **Article 9(2)** continuous-iterative risk-management process throughout the entire lifecycle. The runtime makes the iteration observable as a queue; the Act mandates the iteration as a process.

— **Cognitive Allocation by consequence and uncertainty (Cognition)** maps onto the **risk-based architecture of the Act** itself (Article 6 + Annex III): allocation depth scales with consequence, mirroring the high-risk classification scheme.

These are not paraphrases. They are the same operational topology under different vocabulary and resolution depth.

2. MAPPING TABLE · VEIL V21 RUNTIME × ARTICLE SUBSET

Veil concept	Article(s)	Obligation summary
Execution Kernel (T=0, authority, revocation checks)	**Art. 26§1, §5 · Art. 79(1)**	Conformant use + immediate suspension if §79(1) risk; revocation if revalidation unresolved.
Operator Console (reviewable actions, authority requirements)	**Art. 14(4)(a)–(e)**	Human oversight: understand+monitor (a), automation-bias awareness (b), correct interpretation (c), disregard/override (d), stop/halt (e).
World State + Continuity Kernel (cross-session continuity)	**Art. 26§5 ("reason to consider")**	Pre-revocation notification duty as continuity confidence decays.
Ontology State + Timeline Truth (truth states, constraints)	**Art. 9(2) · Art. 26§5**	Continuous-iterative re-assessment when state diverges.
Contradiction Registry (persistent tensions, pressure)	**Art. 9(2)(b) · Art. 15(4)**	Foreseeable misuse + resilience to errors/inconsistencies + feedback-loop mitigation.
Admitted Memory (only valid memory enters reasoning)	**Art. 9(2)**	Continuous-iterative risk management throughout the entire lifecycle.
Active Recovery (recovered continuity, frames, admission)	**Art. 14 · Art. 26§5**	Human intervention/halt + suspend-resume protocol on §79(1) risk.
Temporal Replay (branch survivability, divergence)	**Art. 12 · Art. 26§6**	Automatic recording + ≥6-month retention under deployer control.
Counterfactuals (no execution authority)	**Art. 9(6)**	Testing for risk-management measures: performance under intended purpose, no execution-state risk.
Branch Arbitration (winning branch, contested branches)	**Art. 25(4) · Art. 26§5**	Written agreement provider ↔ third-party + per-deployer monitoring (see §4).
Meta-Cognition + Reality Model (drift, revalidation)	**Art. 9(2) · Art. 72(1)**	Continuous-iterative process + provider post-market monitoring system.

Cognitive Allocation (depth by consequence)	**Art. 6 · Annex III**	Risk-based classification + deployer obligation subset proportional.

Source: Regulation (EU) 2024/1689, OJ L 12 July 2024. Operational summaries consolidated against Annex A v1.0 §7.1 base cascade (§2 + §5 + §6).

3. THREE DIMENSIONS WHERE VEIL IS MORE GRANULAR THAN THE ACT

— **Risk management.** The Act treats risk management as a single system (Article 9). Veil exposes five distinct pressures (ontology, contradiction, identity, causal, branch). A deployer who maps Article 9(2)'s **"continuous iterative process"** onto Veil receives five independent monitoring signals rather than one composite.

— **Logging.** The Act mandates automatic recording (Article 12) and ≥6-month retention (Article 26§6). Veil maintains two distinct trails: Temporal Replay (executed reality) and Counterfactuals (hypothetical, non-executed). The counterfactual trail evidences Article 9(6) testing without exposing real execution.

— **Human oversight.** The Act mandates oversight (Article 14). Veil maintains two layers: Operator Console (direct human review) and Branch Arbitration (competing trajectories). The branch layer has no direct Article 14 equivalent; it extends oversight to multi-branch resolution.

4. MULTI-DEPLOYER COHERENCE · THE LOAD-BEARING EDGE

The closest article-level anchors are **Article 25(4)** and **Article 26§5**. Article 25(4) requires a written agreement between the provider and third-party suppliers covering information, capabilities, and technical access. Article 26§5 requires each deployer to monitor and notify the provider or distributor and the market-surveillance authority where §79(1) risk emerges. Together they cover the contractual layer (Article 25(4)) and the per-deployer notification layer (Article 26§5).

What remains unnamed at article level is the **runtime synchronisation layer**: when multiple deployers participate in the same evolving consequence environment (**independently coherent institutions drifting into divergent operational realities**), Article 25(4) and Article 26§5 address moments (contract; suspect-risk event), not continuous coherence between concurrent runtime states. The healthcare continuity-fragmentation scenario is the natural pressure-test: a longitudinal deployment where oncology provider, insurer, and research network each carry independent Article 26 obligations under Article 6(1) MDR/IVDR pathway, sharing patient records and clinical decisions across boundaries the regulation does not yet require to synchronise. **This is not a single-system failure but a multi-deployer coherence failure**. The architecture is in the Act; the operational framework is not yet there. A corollary load-bearing thread: the **clock function**. The Act distributes cadence across deployer obligations (Article 26§5, §6), authority review (Article 79), and corrective-action timing (Article 73); continuous recomputation must remain synchronised across all three.

5. LIMITS OF THIS READING

External Reviewer access exposes the runtime/operational layer of the system (Operator, Governance, Continuity, Replay, Branches, Pressure, Cognition, Runtime, Tools). The administrative layer (Audit Logs, System configuration, Overview aggregate) remains under the platform administrator's responsibility and is not assessed here. The mapping reads against observable runtime behaviour, not against administrative governance which remains within the deployer's or operator's own scope under Article 26 and Article 14.

Combined, the runtime layer answers **"is execution still legitimate at the exact moment consequence binds?"** The article-by-article layer answers **"what regulatory consequence attaches to that legitimacy state, article by article?"** The sector-specific stress test the May 9 piece left open, where the bridge survives or breaks, is the healthcare continuity-fragmentation scenario you named on Saturday. When that scenario architecture is ready, the regulatory layer maps onto it through Article 6(1) (MDR/IVDR safety-component pathway), Article 26 paragraph subset specific to clinical deployers (consolidated in Annex A v1.0 §5.4), and Article 25(1)(b) for fine-tuning on patient data. The bridge holds for both directions.

Sprinkling Act · Pre-conformity layer · Brussels · BCE BE 1034.962.482 · sprinklingact.com · References: Shucrani, L. B. (2026). EU AI Act Readiness Report (April 2026), Zenodo DOI 10.5281/zenodo.19671329 · Annex A v1.0 (May 2026), Zenodo DOI 10.5281/zenodo.20042175 · Editorial reading note, non-contractual, not legal advice. Article citations drawn from Regulation (EU) 2024/1689, OJ L 12 July 2024.